

Demo Video Script

Demo video script

A narrated screen-recording script for a ~2.5-minute walkthrough of the project — the kind of thing that works well attached to a portfolio, LinkedIn post, or linked from your resume. Read this once, then record in one or two takes; it's short enough that you don't need to be perfect.

Before you record

1. **Start the app.** Either:

```
docker run --rm -p 8000:8000 --env-file agent/.env wc-bill-review-agent
```

(shows off the containerization — recommended for the video), or the plain `uvicorn api.main:app --reload` route. Either way, confirm `http://127.0.0.1:8000` loads before you hit record.

2. **Tidy the browser window.** Close other tabs, hide the bookmarks bar (`Cmd+Shift+B` in Chrome), and resize the window to something clean — roughly 1280×800 reads well in a recording. Zoom the page to 100%.
3. **Have your bill IDs ready** (from `docs/DEMO.md`) so you're not hunting for them on camera:
 - `WC-BILL-003634` — unbundling
 - `WC-BILL-000066` — clean
 - `WC-BILL-003826` — fragmented billing
4. **Recording tool:** macOS has this built in — press `Cmd+Shift+5` , choose “Record Selected Portion,” drag a box around just the browser window, and make sure your microphone is selected in the “Options” menu before you start.
5. **Do a 10-second mic check** first — record a couple seconds, play it back, confirm your voice isn't clipping or too quiet.

Scene-by-scene script

Timings are approximate — talk at a natural pace, don't rush to hit them exactly. Total target: ~2:30.

[0:00-0:15] Open on the empty app.

(Screen: `http://127.0.0.1:8000` , nothing entered yet)

“This is a bill review assistant I built — it looks at workers' comp medical bills and flags likely coding problems: unbundling, fragmented billing, diagnosis-procedure mismatches, and up-coding. Everything you're about to see runs on 100% synthetic data — no real patients, providers, or claims anywhere in this project — but the coding rules it checks against are real and public: CPT/HCPCS, ICD-10-CM, NCCI edit pairs, and CMS's MUE tables.”

[0:15-0:45] Look up a flagged bill.

(Type `WC-BILL-003634` into the box, press enter, let it load)

“Let's look up a bill. This one trips an unbundling rule — it bills two CPT codes together that, under NCCI's PTP edits, shouldn't be billed together on the same day without a modifier.”

(Point out, as they appear: the “1 issue found” header, the flagged violation row, the CPT code chip)

“The model flags it, and right below, you can see exactly which CPT codes on this specific bill triggered it — that part's computed for free, locally, by re-running the same deterministic rule checker the training data was labeled with, not a second model call.”

[0:45-1:05] Expand the bill details.

(Click “View bill details”)

“Every result also has the actual bill right there if you want to check it yourself — line items, CPT codes, units, charges, diagnosis code, total billed. No need to jump to a separate screen.”

[1:05-1:20] Show a clean bill for contrast.

(Type `WC-BILL-000066` , submit)

“And for contrast — most bills are clean. This one has no issues, so the result is just that: one line, nothing to dig through.”

[1:20-2:00] Turn on AI-generated explanations.

(Check “Generate explanations for flagged violations”, look up `WC-BILL-003826` , wait for it to load)

“There’s one more layer. If I turn on explanations and look up a flagged bill — this one’s a fragmented-billing case, a code billed way more times in one day than its allowed cap — it goes a step further: it retrieves the exact rule text this bill actually violates from a local vector store, and asks an LLM to turn that into a plain-language explanation grounded in this specific bill’s numbers.”

(Point at the longer explanation text once it appears)

“This is the only part of the whole system with a real per-call cost, which is why it’s opt-in — off by default, and I made sure it never fires automatically just because a bill got flagged.”

[2:00-2:20] Quick technical summary.

(Optional: briefly show Docker Desktop with the container running, or just stay on the app)

“Under the hood: the corpus and labels are generated by a rule engine I wrote against those real coding-rule datasets, an XGBoost model does the actual multi-label classification, and the explanation layer is a small RAG pipeline — local embeddings, Azure OpenAI for generation. It’s all wrapped in a FastAPI service and packaged as a Docker image, so it runs the same way anywhere.”

[2:20-2:30] Close.

“That’s the demo — thanks for watching.”

After recording

- **Trim** the start/end in QuickTime Player (`Edit → Trim` , or `Cmd+T`) to cut dead air before/after your narration.
- **Export**: `File → Export As → 1080p` is plenty for a portfolio clip.
- Consider a **captioned/silent cutdown** too (30–45s, just the highlights: lookup → flagged result → explanation) for places that autoplay without sound, like a LinkedIn feed.
- If you want a GIF version for the GitHub README instead of/alongside the video, say so — that’s a different, fully automated route (I can drive an automated click-through and export a looping GIF directly), but it needs the Claude in Chrome browser extension connected first, which isn’t set up in this session yet.